

AIDS Assignment 1

8. 04/05

Q1 What is AI?

→ Artificial Intelligence refers to the simulation of human intelligence in machines. AI systems can perform tasks such as learning, reasoning, problem solving, perception and language understanding.

AI is categorised into various types - including narrow & general AI

Narrow AI : specialised for specific tasks
example :- Siri, Google Assistant

General AI : Hypothetical AI capable of human-like reasoning across domains

Q2 What are AI agents? Explain with example.

→ An AI agent is an entity that perceives its environment through sensors & acts upon its perception action cycle to achieve a goal.

AI agents can be classified as:

- (1) Simple Reflex Agents : Reacts to current percepts
- (2) Model-Based Agents : Maintain internal model

Q2

Q'

- (i) Goal Based Agents : Maintains internal models
- (ii) Utility based : Maximizes utility functions

The example of each of the different agents is that :

Simple Reflex Agents (Rule-Based)

Thermostat :- Turns on/off heat based on temperature

Model-Based Reflex Agents

Self driving cars : Use sensor data to track obstacles & road conditions

Utility-Based Agents

Stock marketing bots : Optimise for maximising profit while minimising risk

Learning Agent :

Recommendation System :- Learn from preference

Q1 AI plays a crucial role in various aspects of healthcare & daily life, such as:

- Early detection & diagnosis : AI-based models detected and analyzed patterns
- Medical Research & Drug Discovery : It discovers about potential treatment
- Healthcare automation, Supply chain automation & remote sensing are few other ways AI helped during the pandemic

Q3 How AI is used to solve 8 puzzle problem?

→ The 8-puzzle problem is a sliding puzzle consisting of ~~3x3~~ grid with eight tiles and one empty space. AI techniques used to solve it include:

- Breadth First search: explores all possible moves level by level
- Depth First search: explores all deeper path first
- A* Algorithm: Uses a heuristic function to find optimal path

Nikita Chhabra

Roll no. 07

D15C

04
05

subject, syntactic
preference, genre data

Autolander

heading precision, passenger safety
Runway wind, attitude
Flaps, engines, landing gear
Altimeter, gyroscope, GPS

Sentry dog

weak detection, accuracy
visible area, shoulder
in movement

camera, motion sensors

lapping bot for an offline Bookstar

ability : Partially Observable
vision : Stochastic

: Sequential

Dynamic : Dynamic

continuous : discrete

Multi-Agent : Multi Agent

- Greedy Best First Search :- Priorities moves based on heuristics without considering past cost

Q4 The PEAS framework stands for Performance measure, Environment, Actuators & Sensors.
It is used to define the components of an AI agent by specifying how it interacts with its environment & evaluates its success.

- Taxi Driver

~~Performance :- Safety, speed, passenger comfort~~

~~Environment :- Roads, passengers, traffic~~

~~Actuator :- steering, acceleration, brakes~~

~~Sensors :- GPS, cameras, speedometer~~

- Medical Diagnosis System

P :- Diagnosis accuracy, treatment success

E :- Patients, symptoms, diseases

A :- Recommendations, prescription

S :- Patient history, lab tests

- Music Composer A I

P :- Harmony, originality

E :- Musical notes, instruments

A :- MIPI output , synthesize

S :- user preference , genre data

- Aircraft Autolander

P :- landing precision , passenger safety

E :- Runway wind , altitude

A :- Flaps , engines , landing gear

S :- Altimeter , gyroscope , GPS

- Robotic Sentry Iyer

P :- Threat detection , accuracy

E :- Security area , border

A :- Iyer movement

S :- cameras , motion sensors

Q5 A shopping bot for an offline Bookstore

Observability : Partially observable

Determinism : Stochastic

Episodic : Sequential

Static / Dynamic : Dynamic

Discrete / continuous : Discrete

Singer / Multi-Agent : Multi-Agent

96

Feature

Model Based Agent

Utility Based

Internal Model

Maintains an internal model of the world

evaluates actions based on utility

Decision Making

Uses past state to make decisions

Selects action to maximize utility

Example

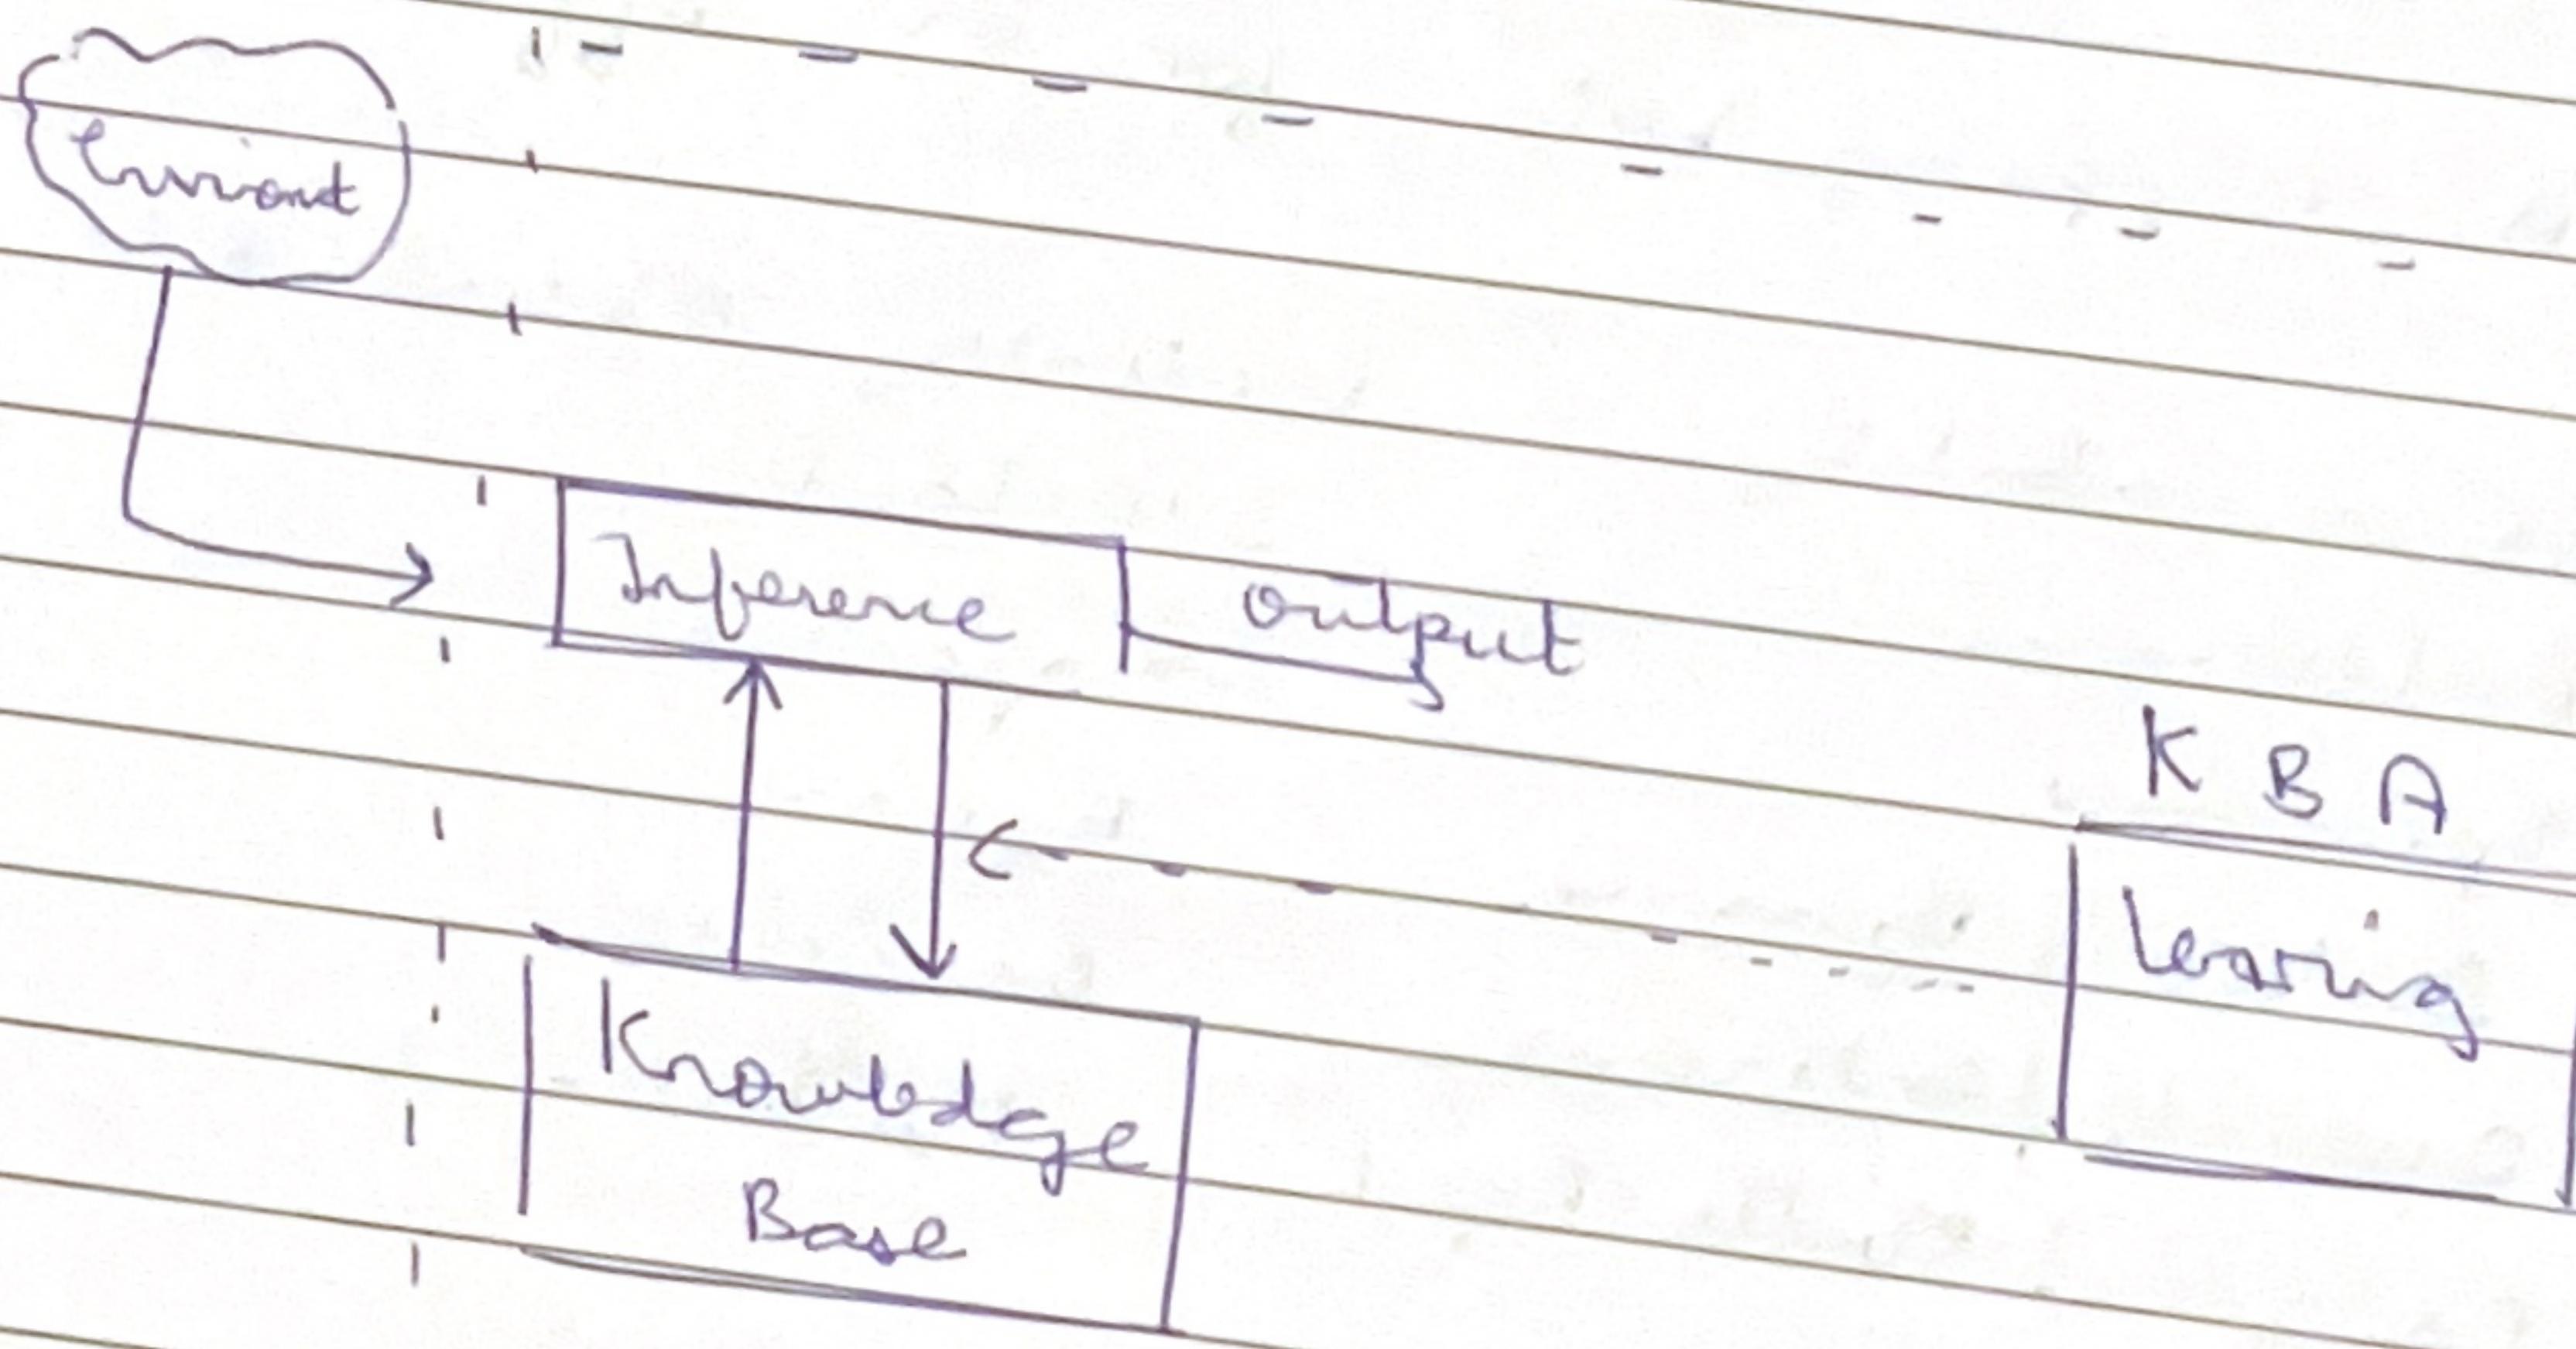
Self driving cars
car Navigating through traffic

AJ recommending personalized products

Q7

Architecture
Agents

of knowledge based & learning



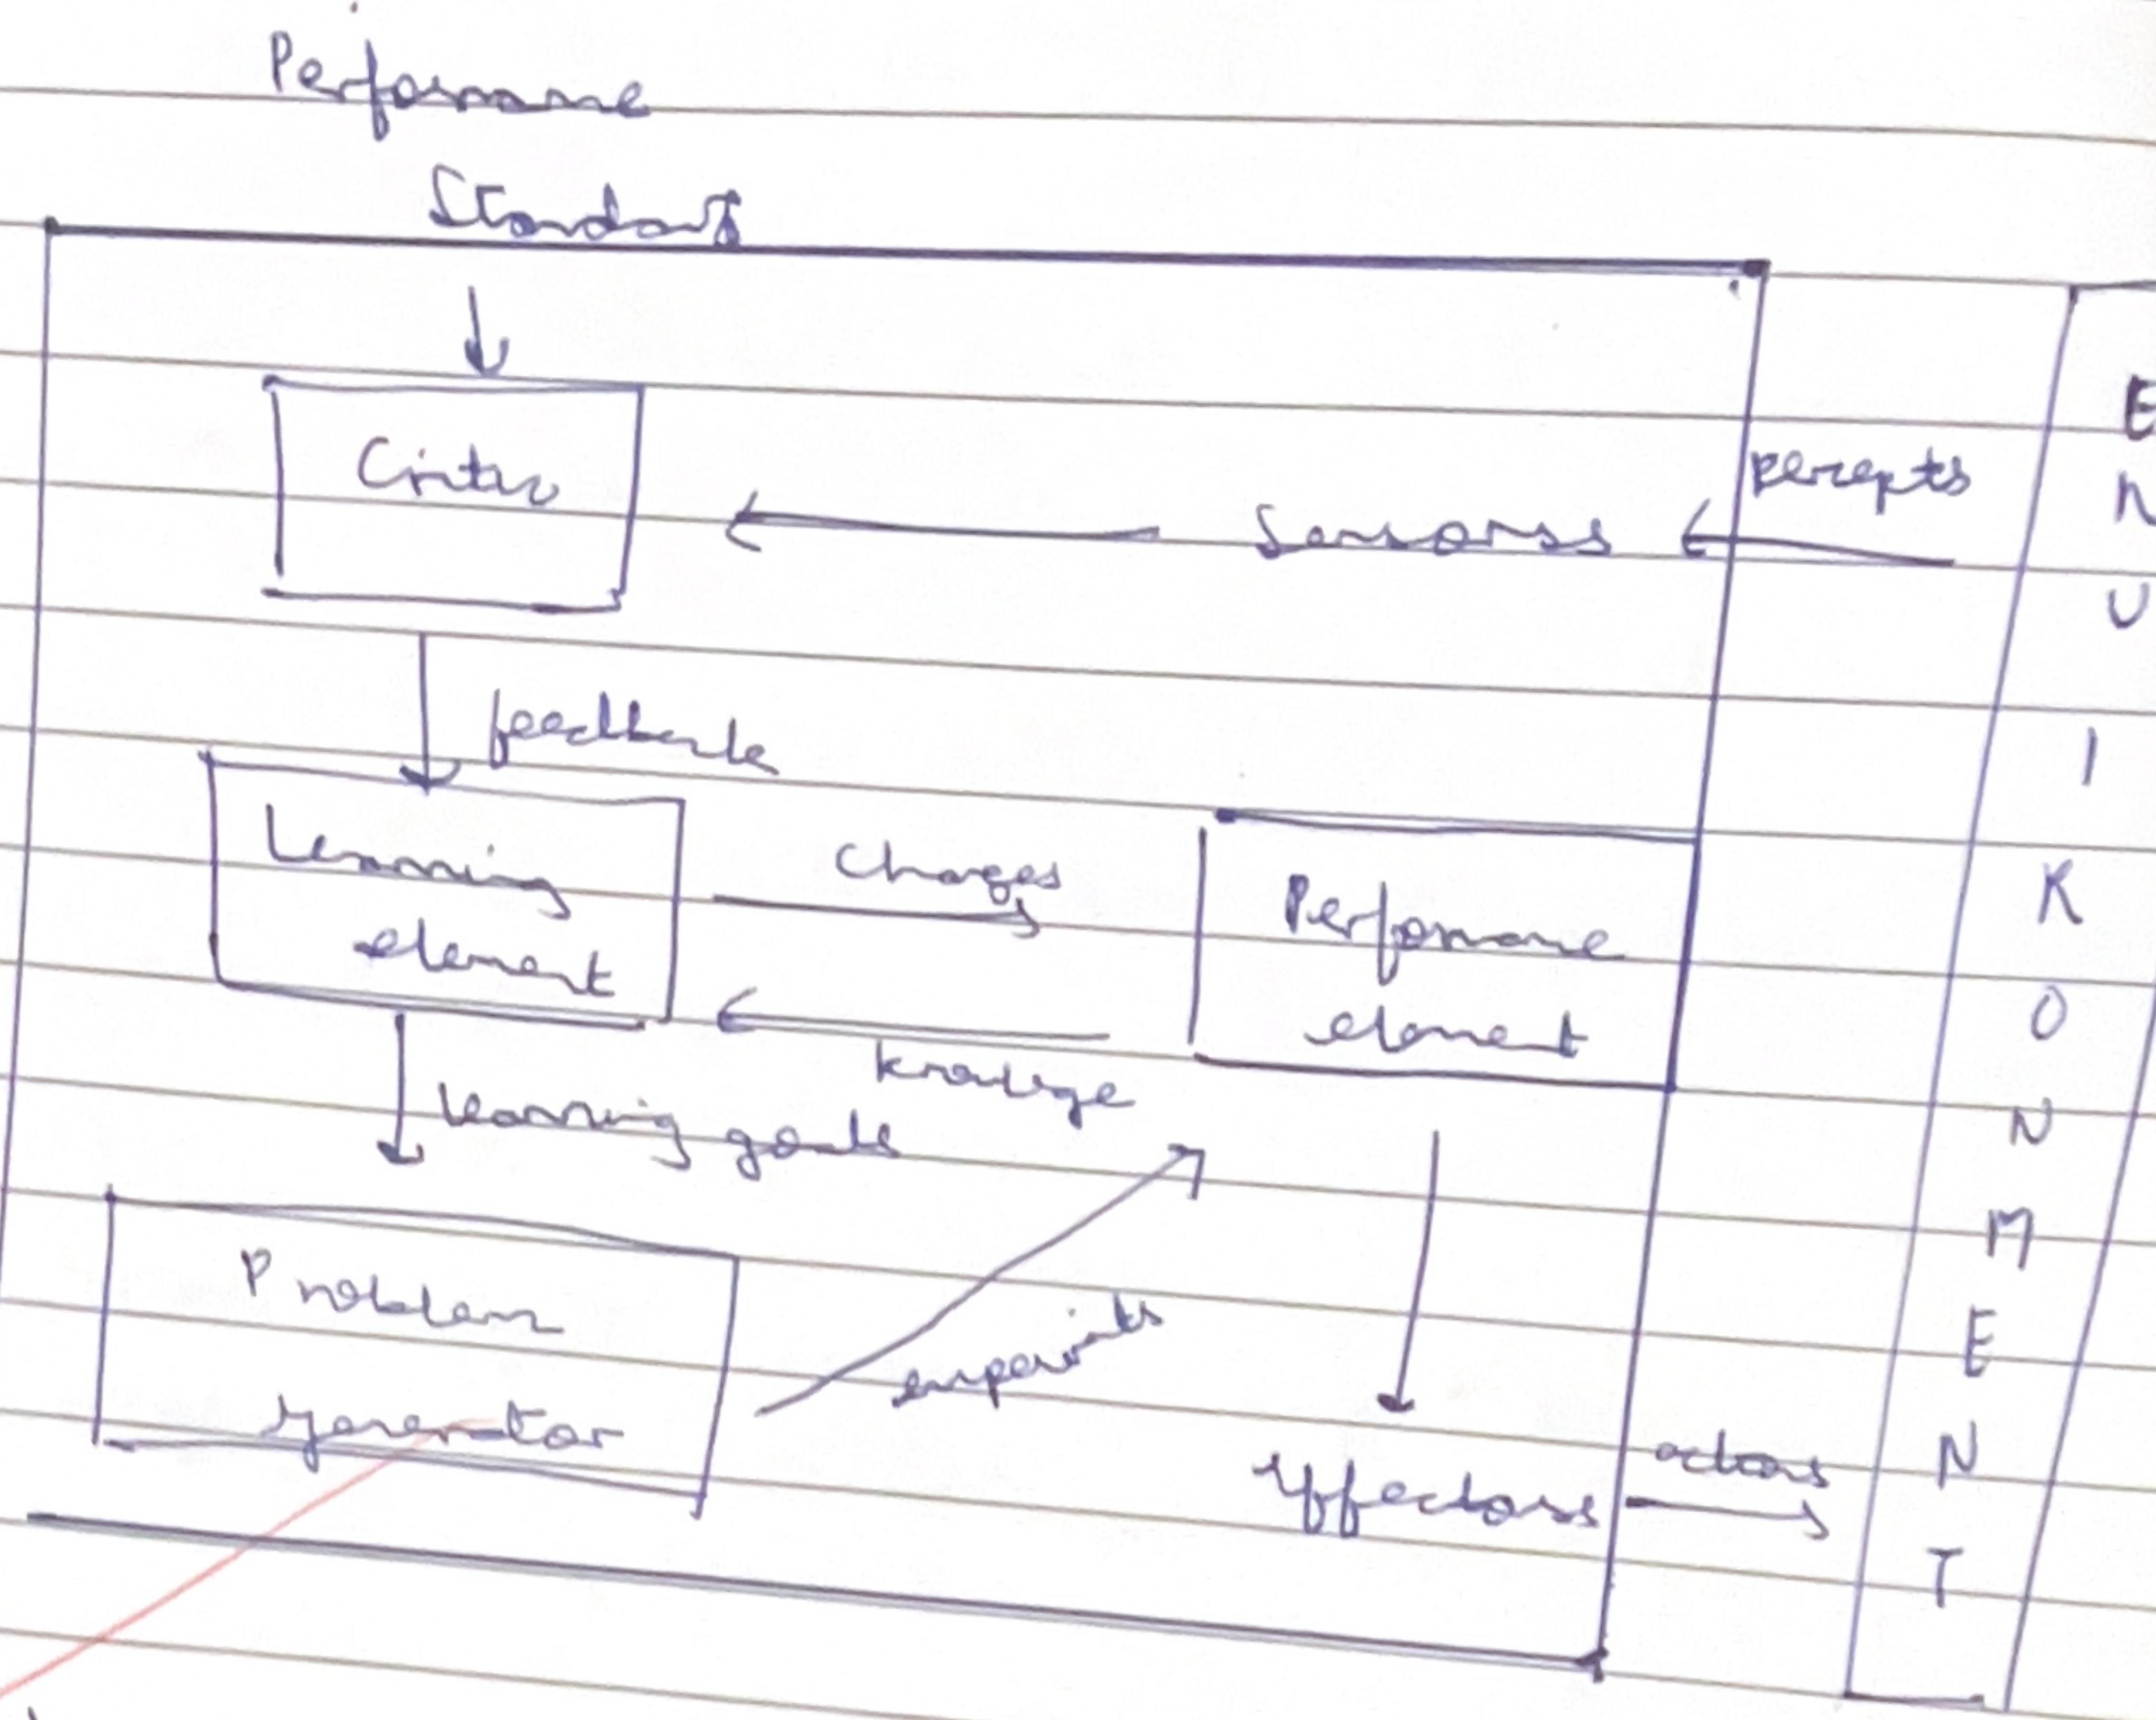
A Knowledge Based Agent uses stored knowledge to make decisions and reason about the environment. It's architecture consists of

- ① Knowledge Base : - Stores knowledge to make decisions and reason about the environment. Its architecture consists of :-
 - ① Knowledge Base (KB) : Stores facts, rules and reason about in a structured format
 - ② Inference Engine : Applies logical reasoning to derive conclusion from the Knowledge base
 - ③ Perception Module : Gather inputs from sensors or external sources
 - ④ Action Module : Invokes action based on derived conclusions
 - ⑤ Learning Mechanism : Updates the knowledge base with new information

Example :- A medical diagnosis system that stores diseases symptoms & applies logical rules

Q 7

Architecture of Learning Agent



A learning Agent improves its performance over time by learning from past experiences. It has few main components.

- Learning Agent : Learns from interactions & updates knowledge
- Performance Element : Uses the learned knowledge to make decisions
- Criteria : Evaluates the agent's performance by comparing actions with desired outcomes

Predicate logic conversion

(a) Anita travels by car if available, otherwise, travels by bus

Available (Car) \Rightarrow Travels (Anita, Car)

\neg Available (Car) \rightarrow Travels (Anita, Bus)

(b) Bus goes via Andheri & Goregaon

Routes (Bus, Andheri)

Route (Bus, Goregaon)

(c) Car has a puncture, so it is not available

Puncture (Car) $\Rightarrow \neg$ Available (Car)

Forward Reasoning

From (3) \Rightarrow Available (Car)

(1) Has Puncture (Car) $\rightarrow \neg$ Car Available (Car)

(2) \neg Car Available (Car) \rightarrow Travels (Anita, Bus)

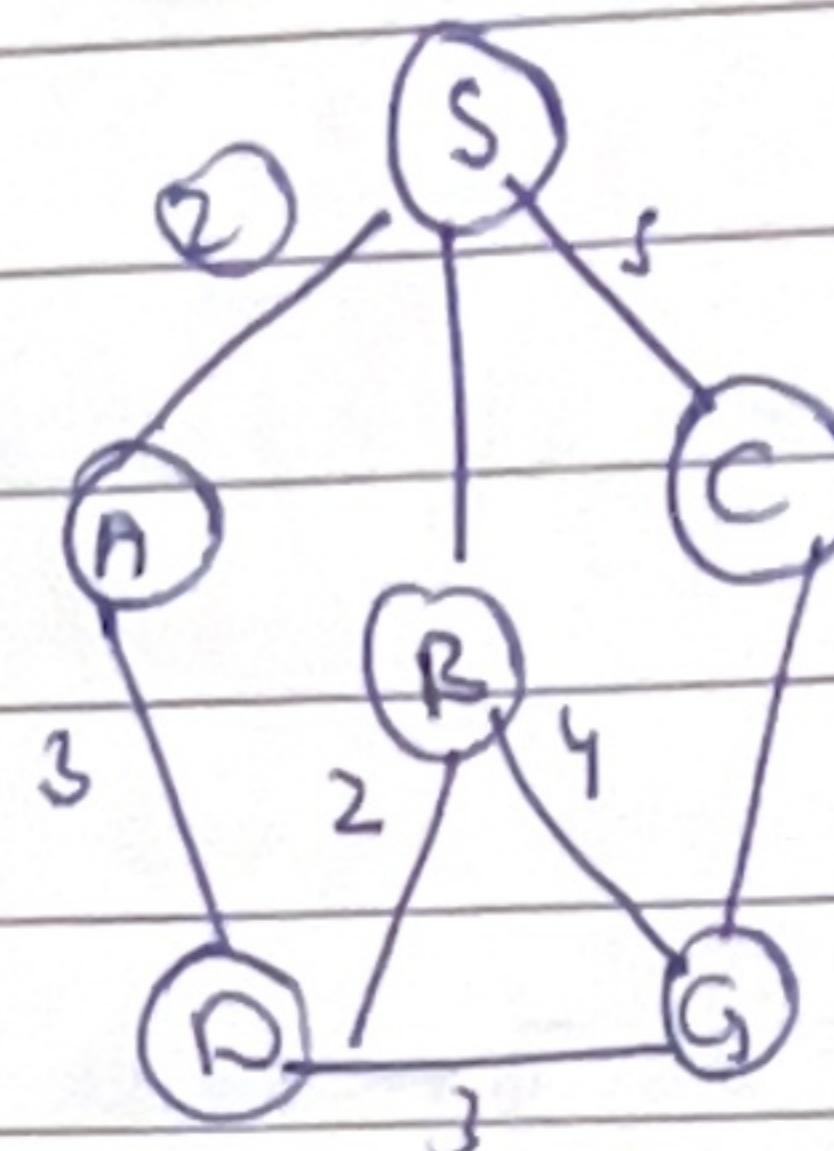
(3) Travels (Anita, Bus) [from 1 & 2]

(4) Bus goes via Goregaon [from b]

(5) Travels (Anita, Goregaon)

Yes, Anita will travel via Goregaon.

Q9



$S \rightarrow D$ cost(2)

$S \rightarrow B$ cost 5

$S \rightarrow C$ cost 5

$A \rightarrow D$ cost 3

$A \rightarrow G$ (cost 3)

$C \rightarrow G$ (cost 4)

BFS Traversal from S to G

- ① Start at S \rightarrow Queue [S]
- ② Expand S \rightarrow Queue [A, B, C]
- ③ Expand B (no new nodes) \rightarrow Queue [C, G]
- ④ Expand C \rightarrow Queue [D, G]
- ⑤ Expand D \rightarrow Queue [G]
- ⑥ Expand G \rightarrow Goal Found

Shortest path S to G is:

$S \rightarrow A \rightarrow G$

Total Cost : $2 + 3 = 5$

Q10 Iterative Deepening Search (IDS)

It combines the space efficiency of DFS & completeness of BFS by repeatedly running DLS with increasing depth limit

Advantage: guarantees finding the shortest path while using less memory than BFS

Q12 Hill Climbing Algorithm is a local search algorithm that continuously moves towards the best neighboring state with a higher heuristic value, aiming to reach a global optimum

Steps:

- (1) Start from an initial state
- (2) Evaluate neighboring states and move to the one with the highest heuristic value
- (3) Repeat until no better neighbor

Trap example:-

Imagine a mountain climbing scenario where a riker moves uphill based on the steepest slope. If they reach a peak that is not the highest (global maximum), they might get stuck.

Q12 Drawbacks of Hill Climbing

- Local Minimum

The algorithm may stop at peak that is not the global optimum

- Plateau Problem:

A flat region where all neighbouring states have the same heuristic value, causing search to halt

- No Back tracking

Once it moves forward, it cannot recover from a bad decision

= Limitation of Steepest - Ascent Hill Climbing

~~Steepest - Ascent Hill Climbing~~ selects the best possible move in each step but has extra limitation

- ① Slow progress in Narrow ridges
- ② More likely to get stuck in local minima
- ③ Inefficienc in large search space

Solutions :

Random Restart

Simulated Annealing

Simulated Annealing is a probabilistic search technique that avoids getting trapped in local optima by allowing occasional bad moves

Steps :

- 1) Start with random solution
- 2) Evaluate its cost
- 3) Move to a neighboring solution
- 4) Accept it if it better, otherwise, accept it with probability decreasing over time
- 5) Reduce "temperature" gradually
- c) Repeat until convergence

Algorithm :

Initiate temperature T , cooling rate α .

Choose an initial solution s

Repeat until $T \rightarrow 0$:

- a. Select a neighboring solution s'
- b. Compute cost diff $\Delta E = \text{cost}(s) - \text{cost}(s')$
- c. If $\Delta E < 0$, move to s'
- d. Else, move to s' with prob $e^{\Delta E/T}$
- e. Reduce T : $T = \alpha T$

Returns best solution found

In Travelling Salesman Problem where a delivery truck must find shortest route

f 14

A* is an informed search Algorithm that uses

$$f(n) = g(n) + h(n)$$

where $g(n)$ is the cost from start to node n & $h(n)$ is heuristic

e.g. Finding shortest path in a graph like Google maps

It is widely used in pathfinding, AI & game development

Uses both $g(n)$ & $h(n)$

A* is optimal

A* is complete

f 15

The Minimax Algorithm is used in adversarial search (e.g. two-player games) to determine the optimal move by assuming both players play optimally

Maximizer (e.g. X) tries to get the highest score
Minimizer (e.g. O) tries to reduce the score

Algorithm

- ① Generate the game tree up to depth limit
- ② Assign heuristic values to leaf nodes
- ③ Backpropagate values:
 - Maximizer picks the maximum value

• Minimax picks the minimax value

The root node selects the best move based on propagated

game tree: A minimax tree for Tic-Tac-Toe would show all possible board states, evaluating the best move at each step

Alpha-Beta Pruning optimizes minimax by skipping unnecessary branches, reducing computations. It uses two values

~~Alpha (α): Best score minimizer can achieve~~

~~Beta (β): Best score maximizer can achieve~~

Algorithm:

1) Perform Minimax search

2) Prune branches where:

if Minimax best (α) \geq Minimax best (β),

further evaluation is stopped

if Minimax's best (β) \leq Minimax's best (α)

3) This reduces time complexity without affecting the final decision

Example: In a game tree, if a branch has already provided a worse outcome than the best known

Q16 more, it is ignored to save time

Q17

The 'Wumpus' world is a grid-based PIZ environment where an agent explores a cave while avoiding hazards like pits & the wumpus monster

PEAS descriptor

Performance Measure: Reaching the gold safely, minimizing steps

Environment: A 4×4 grid with pits, gold & wumpus

Actuators: Move Forward, turn, grab, shoot, climb

Sensors: Perceive stench, breeze, glitter

Percept sequence generation:

- 1) Agent starts at $(1, 1)$, sensing its surroundings
- 2) If stench is detected, the Wumpus is nearby
- 3) If breeze is detected, a pit is nearby
- 4) The agent infers safe paths & navigates towards the gold while avoiding hazards

$$\begin{array}{r}
 75N0 \\
 10R5 \\
 \hline
 10NS2
 \end{array}$$

As $E = 5$ the sum of NK should
be a carry.

$$\therefore N = 6, R = 8$$

$$\begin{array}{r}
 9560 \\
 1085 \\
 \hline
 10654
 \end{array}$$

Assigning the remain value to D

$$\therefore D = 7$$

$$\begin{array}{r}
 9567 \\
 1085 \\
 \hline
 10652
 \end{array}$$

Solution:-

$$S = 9$$

$$M = 1$$

$$E = 5$$

$$D = 0$$

$$N = 6$$

$$N = 6$$

$$D = 7$$

$$G = 5$$

$$Y = 2$$

Q19 Modus Ponens is a fundamental rule of inference in propositional logic

If $P \rightarrow Q$, P is true, then Q is also true

Example :- If it rains, the ground will be wet ($P \rightarrow Q$)

1) It is raining (P is true)

3) Therefore, the ground is wet (Q is true)
mathematically,

$$\begin{array}{c} P \rightarrow Q \\ P \\ \therefore Q \end{array}$$

- Q20 ① Forward Chaining is rule driven
② it starts from known facts & applies inference rules to derive new facts until the goal is reached
③ it is used in expert systems

Example :-

Fat - Sore Throat

Rule :- If sore throat \rightarrow infection

New Fact :- Infection

Rule :- If infection \rightarrow need antibiotic

Conclusion :- "Needs to Antibiotic"

Q20 Backward Chaining

- (1) It is more goal driven inference.
- (2) It starts from goal & works backwards to determine facts needed to prove it.
- (3) It is used in prolog programming.

Example:-

goal : Does the patient need antibiotics?

- (1) Clerk : Does the patient have an infection?
- (2) Clerk : Does the patient have a sore throat?
- (3) If both hold, conclude " Need antibiotics"

This reduces unnecessary computation by only exploring relevant facts

J: